

```
In [5]: # IMPORT LIBRARIES

import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.layers import (
    Input, Conv2D, MaxPooling2D, Flatten,
    Dense, Dropout, GlobalAveragePooling2D
)
from tensorflow.keras.applications import (
    VGG16, ResNet50, MobileNet, InceptionV3, EfficientNetB0
)
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.metrics import BinaryAccuracy
from tensorflow.keras.models import load_model
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

```
In [6]: # LOAD THE DATA

import os

base_dir = r"D:\DS PROJECTS SUBMISSION\Multiclass_Fish_Image_Classification\Dataset"

print("Inside Dataset folder:")
print(os.listdir(base_dir))

print("Dataset extracted successfully!")

Inside Dataset folder:
['test', 'train', 'val']
Dataset extracted successfully!
```

```
In [7]: # DATA RESCALING

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=30,
    zoom_range=0.2,
    horizontal_flip=True
)

val_datagen = ImageDataGenerator(
```

```
    rescale=1./255
)
```

In [8]: # GENERATORS

```
train_generator = train_datagen.flow_from_directory(
    r"D:\WORKOUTS\DATA_CLEANING\Dataset CSV\MINI PROJECT 5\Dataset\images.cv_jzk611",
    target_size=(224,224),
    batch_size=32,
    class_mode='categorical'
)

val_generator = val_datagen.flow_from_directory(
    r"D:\WORKOUTS\DATA_CLEANING\Dataset CSV\MINI PROJECT 5\Dataset\images.cv_jzk611",
    target_size=(224,224),
    batch_size=32,
    class_mode='categorical'
)

print(train_generator.samples)
print(val_generator.samples)

NUM_CLASSES = train_generator.num_classes
print("Classes:", train_generator.class_indices)
```

Found 6225 images belonging to 11 classes.

Found 1092 images belonging to 11 classes.

6225

1092

Classes: {'animal fish': 0, 'animal fish bass': 1, 'fish sea_food black_sea_sprat': 2, 'fish sea_food gilt_head_bream': 3, 'fish sea_food hourse_mackerel': 4, 'fish sea_food red_mullet': 5, 'fish sea_food red_sea_bream': 6, 'fish sea_food sea_bass': 7, 'fish sea_food shrimp': 8, 'fish sea_food striped_red_mullet': 9, 'fish sea_food trout': 10}

In [14]: # TRAIN THE MODEL

```
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(224,224,3)),
    MaxPooling2D(2,2),

    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(2,2),

    Conv2D(128, (3,3), activation='relu'),
    MaxPooling2D(2,2),

    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),

    Dense(NUM_CLASSES, activation='softmax')
])

model.compile(
    optimizer=Adam(learning_rate=0.001),
```

```

        loss="categorical_crossentropy",
        metrics=['accuracy']
    )

model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_97 (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d_7 (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_98 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_8 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_99 (Conv2D)	(None, 52, 52, 128)	73,856
max_pooling2d_9 (MaxPooling2D)	(None, 26, 26, 128)	0
flatten_1 (Flatten)	(None, 86528)	0
dense_12 (Dense)	(None, 128)	11,075,712
dropout_6 (Dropout)	(None, 128)	0
dense_13 (Dense)	(None, 11)	1,419

Total params: 11,170,379 (42.61 MB)

Trainable params: 11,170,379 (42.61 MB)

Non-trainable params: 0 (0.00 B)

In [15]: # FUNCTION FOR PRETRAINED MODELS

```

def build_pretrained_model(base_model, num_classes):
    base_model.trainable = False    # Freeze pretrained weights

    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(256, activation='relu')(x)
    x = Dropout(0.5)(x)

    # ✓ MULTI-CLASS OUTPUT
    output = Dense(num_classes, activation='softmax')(x)

    model = Model(inputs=base_model.input, outputs=output)

    model.compile(
        optimizer=Adam(learning_rate=0.001),
        loss='categorical_crossentropy',
        metrics=['accuracy']
    )

```

```
    return model
```

```
In [16]: # LOAD TRAINED MODELS (COMPARISON)

vgg_model = build_pretrained_model(
    VGG16(weights='imagenet', include_top=False, input_shape=(224,224,3)),
    NUM_CLASSES
)

resnet_model = build_pretrained_model(
    ResNet50(weights='imagenet', include_top=False, input_shape=(224,224,3)),
    NUM_CLASSES
)

mobilenet_model = build_pretrained_model(
    MobileNet(weights='imagenet', include_top=False, input_shape=(224,224,3)),
    NUM_CLASSES
)

inception_model = build_pretrained_model(
    InceptionV3(weights='imagenet', include_top=False, input_shape=(224,224,3)),
    NUM_CLASSES
)

efficientnet_model = build_pretrained_model(
    EfficientNetB0(weights='imagenet', include_top=False, input_shape=(224,224,3)),
    NUM_CLASSES
)
```

```
In [17]: # SAVE BEST MODEL

checkpoint = ModelCheckpoint(
    "best_model.h5",
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
    verbose=1
)

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True,
    verbose=1
)
```

```
In [18]: # TRAIN THE BEST MODEL

# EfficientNet
checkpoint_eff = ModelCheckpoint(
    "best_efficientnet.h5",
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
```

```
    verbose=1
)

history_efficientnet = efficientnet_model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=3,
    callbacks=[checkpoint_eff, early_stop]
)

# VGG16
checkpoint_vgg = ModelCheckpoint(
    "best_vgg16.h5",
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
    verbose=1
)

history_vgg = vgg_model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=3,
    callbacks=[checkpoint_vgg, early_stop]
)

# ResNet50
checkpoint_resnet = ModelCheckpoint(
    "best_resnet50.h5",
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
    verbose=1
)

history_resnet = resnet_model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=3,
    callbacks=[checkpoint_resnet, early_stop]
)

# MobileNet
checkpoint_mobilenet = ModelCheckpoint(
    "best_mobilenet.h5",
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
    verbose=1
)

history_mobilenet = mobilenet_model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=3,
```

```

        callbacks=[checkpoint_mobilenet, early_stop]
    )

# InceptionV3
checkpoint_inception = ModelCheckpoint(
    "best_inception.h5",
    monitor='val_accuracy',
    save_best_only=True,
    mode='max',
    verbose=1
)

history_inception = inception_model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=3,
    callbacks=[checkpoint_inception, early_stop]
)

print("Training complete!")

```

Epoch 1/3
195/195 0s 2s/step - accuracy: 0.1571 - loss: 2.3647
Epoch 1: val_accuracy improved from None to 0.17125, saving model to best_efficientnet.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Epoch 1: finished saving model to best_efficientnet.h5
195/195 496s 2s/step - accuracy: 0.1611 - loss: 2.3430 - val_accuracy: 0.1712 - val_loss: 2.3174

Epoch 2/3
195/195 0s 2s/step - accuracy: 0.1568 - loss: 2.3204
Epoch 2: val_accuracy did not improve from 0.17125
195/195 345s 2s/step - accuracy: 0.1692 - loss: 2.3185 - val_accuracy: 0.1712 - val_loss: 2.3101

Epoch 3/3
195/195 0s 2s/step - accuracy: 0.1793 - loss: 2.3072
Epoch 3: val_accuracy did not improve from 0.17125
195/195 346s 2s/step - accuracy: 0.1761 - loss: 2.3077 - val_accuracy: 0.1712 - val_loss: 2.3131

Restoring model weights from the end of the best epoch: 2.

Epoch 1/3
195/195 0s 7s/step - accuracy: 0.3329 - loss: 1.9826
Epoch 1: val_accuracy improved from None to 0.73352, saving model to best_vgg16.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
Epoch 1: finished saving model to best_vgg16.h5
195/195 1771s 9s/step - accuracy: 0.4683 - loss: 1.6390 - val_accuracy: 0.7335 - val_loss: 1.0808
Epoch 2/3
195/195 0s 10s/step - accuracy: 0.6749 - loss: 1.0552
Epoch 2: val_accuracy improved from 0.73352 to 0.87454, saving model to best_vgg16.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
Epoch 2: finished saving model to best_vgg16.h5
195/195 2248s 12s/step - accuracy: 0.7033 - loss: 0.9673 - val_accuracy: 0.8745 - val_loss: 0.6770
Epoch 3/3
195/195 0s 9s/step - accuracy: 0.7633 - loss: 0.7736
Epoch 3: val_accuracy improved from 0.87454 to 0.89286, saving model to best_vgg16.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
Epoch 3: finished saving model to best_vgg16.h5
195/195 2147s 11s/step - accuracy: 0.7849 - loss: 0.7110 - val_accuracy: 0.8929 - val_loss: 0.4905
Restoring model weights from the end of the best epoch: 3.
Epoch 1/3
195/195 0s 4s/step - accuracy: 0.1572 - loss: 2.4052
Epoch 1: val_accuracy improved from None to 0.17125, saving model to best_resnet50.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
Epoch 1: finished saving model to best_resnet50.h5
195/195 812s 4s/step - accuracy: 0.1690 - loss: 2.3202 - val_accuracy: 0.1712 - val_loss: 2.2321
Epoch 2/3
195/195 0s 4s/step - accuracy: 0.1814 - loss: 2.2289
Epoch 2: val_accuracy improved from 0.17125 to 0.19963, saving model to best_resnet50.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
Epoch 2: finished saving model to best_resnet50.h5
195/195 835s 4s/step - accuracy: 0.1846 - loss: 2.2227 - val_accuracy: 0.1996 - val_loss: 2.1801
Epoch 3/3
195/195 0s 3s/step - accuracy: 0.1976 - loss: 2.1739
Epoch 3: val_accuracy improved from 0.19963 to 0.22894, saving model to best_resnet50.h5
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
```

```
Epoch 3: finished saving model to best_resnet50.h5  
195/195 789s 4s/step - accuracy: 0.1973 - loss: 2.1711 - val_accuracy: 0.2289 - val_loss: 2.1019
```

```
Restoring model weights from the end of the best epoch: 1.
```

```
Epoch 1/3
```

```
195/195 0s 1s/step - accuracy: 0.6519 - loss: 1.1021
```

```
Epoch 1: val_accuracy improved from None to 0.97344, saving model to best_mobilenet.h5
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
```

```
Epoch 1: finished saving model to best_mobilenet.h5
```

```
195/195 320s 2s/step - accuracy: 0.8182 - loss: 0.5828 - val_accuracy: 0.9734 - val_loss: 0.0975
```

```
Epoch 2/3
```

```
195/195 0s 1s/step - accuracy: 0.9400 - loss: 0.1840
```

```
Epoch 2: val_accuracy improved from 0.97344 to 0.97894, saving model to best_mobilenet.h5
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
```

```
Epoch 2: finished saving model to best_mobilenet.h5
```

```
195/195 305s 2s/step - accuracy: 0.9489 - loss: 0.1567 - val_accuracy: 0.9789 - val_loss: 0.0605
```

```
Epoch 3/3
```

```
195/195 0s 1s/step - accuracy: 0.9677 - loss: 0.1054
```

```
Epoch 3: val_accuracy improved from 0.97894 to 0.98168, saving model to best_mobilenet.h5
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
```

```
Epoch 3: finished saving model to best_mobilenet.h5
```

```
195/195 303s 2s/step - accuracy: 0.9700 - loss: 0.0988 - val_accuracy: 0.9817 - val_loss: 0.0503
```

```
Restoring model weights from the end of the best epoch: 3.
```

```
Epoch 1/3
```

```
195/195 0s 2s/step - accuracy: 0.6307 - loss: 1.1889
```

```
Epoch 1: val_accuracy improved from None to 0.90018, saving model to best_inception.h5
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
```

```

Epoch 1: finished saving model to best_inception.h5
195/195 544s 3s/step - accuracy: 0.7635 - loss: 0.7407 - val_ac
curacy: 0.9002 - val_loss: 0.2824
Epoch 2/3
195/195 0s 2s/step - accuracy: 0.8674 - loss: 0.4025
Epoch 2: val_accuracy improved from 0.90018 to 0.95421, saving model to best_incepti
on.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.
saving.save_model(model)`. This file format is considered legacy. We recommend using
instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.savin
g.save_model(model, 'my_model.keras')`.
Epoch 2: finished saving model to best_inception.h5
195/195 508s 3s/step - accuracy: 0.8848 - loss: 0.3463 - val_ac
curacy: 0.9542 - val_loss: 0.1571
Epoch 3/3
195/195 0s 2s/step - accuracy: 0.8993 - loss: 0.2905
Epoch 3: val_accuracy improved from 0.95421 to 0.95513, saving model to best_incepti
on.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.
saving.save_model(model)`. This file format is considered legacy. We recommend using
instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.savin
g.save_model(model, 'my_model.keras')`.
Epoch 3: finished saving model to best_inception.h5
195/195 506s 3s/step - accuracy: 0.9049 - loss: 0.2723 - val_ac
curacy: 0.9551 - val_loss: 0.1356
Restoring model weights from the end of the best epoch: 1.
Training complete!

```

In [19]: # BEST MODEL

```

model_results = {
    "EfficientNetB0": max(history_efficientnet.history['val_accuracy']),
    "VGG16": max(history_vgg.history['val_accuracy']),
    "ResNet50": max(history_resnet.history['val_accuracy']),
    "MobileNet": max(history_mobilenet.history['val_accuracy']),
    "InceptionV3": max(history_inception.history['val_accuracy'])
}

best_model_name = max(model_results, key=model_results.get)
best_accuracy = model_results[best_model_name]

print("Best Model:", best_model_name)
print("Best Validation Accuracy:", best_accuracy)

```

Best Model: MobileNet
 Best Validation Accuracy: 0.9816849827766418

In [20]: # CHECK TRAIN VS VAL ACCURACY

```

print("Train Acc:", max(history_mobilenet.history['accuracy']))
print("Val Acc:", max(history_mobilenet.history['val_accuracy']))

```

Train Acc: 0.9699598550796509
 Val Acc: 0.9816849827766418

In [21]: # FINE TUNING

```

mobilenet_model.trainable = True

for layer in mobilenet_model.layers[:-20]:
    layer.trainable = False

mobilenet_model.compile(
    optimizer=Adam(learning_rate=1e-5),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

fine_tune_history = mobilenet_model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=5,
    callbacks=[checkpoint_mobilenet, early_stop]
)

```

Epoch 1/5

195/195 0s 1s/step - accuracy: 0.9716 - loss: 0.0883

Epoch 1: val_accuracy improved from 0.98168 to 0.99176, saving model to best_mobilenet.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Epoch 1: finished saving model to best_mobilenet.h5

195/195 320s 2s/step - accuracy: 0.9725 - loss: 0.0852 - val_accuracy: 0.9918 - val_loss: 0.0335

Epoch 2/5

195/195 0s 1s/step - accuracy: 0.9789 - loss: 0.0723

Epoch 2: val_accuracy improved from 0.99176 to 0.99267, saving model to best_mobilenet.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Epoch 2: finished saving model to best_mobilenet.h5

195/195 320s 2s/step - accuracy: 0.9782 - loss: 0.0700 - val_accuracy: 0.9927 - val_loss: 0.0290

Epoch 3/5

195/195 0s 1s/step - accuracy: 0.9826 - loss: 0.0608

Epoch 3: val_accuracy improved from 0.99267 to 0.99451, saving model to best_mobilenet.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```

Epoch 3: finished saving model to best_mobilenet.h5
195/195 315s 2s/step - accuracy: 0.9825 - loss: 0.0605 - val_ac
curacy: 0.9945 - val_loss: 0.0267
Epoch 4/5
195/195 0s 1s/step - accuracy: 0.9822 - loss: 0.0548
Epoch 4: val_accuracy did not improve from 0.99451
195/195 283s 1s/step - accuracy: 0.9827 - loss: 0.0554 - val_ac
curacy: 0.9945 - val_loss: 0.0261
Epoch 5/5
195/195 0s 1s/step - accuracy: 0.9804 - loss: 0.0589
Epoch 5: val_accuracy improved from 0.99451 to 0.99542, saving model to best_mobilenet.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.
saving.save_model(model)`. This file format is considered legacy. We recommend using
instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.savin
g.save_model(model, 'my_model.keras')`.

Epoch 5: finished saving model to best_mobilenet.h5
195/195 284s 1s/step - accuracy: 0.9815 - loss: 0.0575 - val_ac
curacy: 0.9954 - val_loss: 0.0245
Restoring model weights from the end of the best epoch: 5.

```

In [22]: # BEST MODEL

```
best_model = tf.keras.models.load_model("best_mobilenet.h5")
```

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be buil
t. `model.compile_metrics` will be empty until you train or evaluate the model.

```

In [23]: # GET TRUE LABELS & PREDICTIONS

```

# LOAD BEST MODEL (MobileNet)
best_model = tf.keras.models.load_model("best_mobilenet.h5")

# RESET GENERATOR
val_generator.reset()

# PREDICT
y_pred_prob = best_model.predict(val_generator)
y_pred = np.argmax(y_pred_prob, axis=1)

# TRUE LABELS
y_true = val_generator.classes

```

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be buil
t. `model.compile_metrics` will be empty until you train or evaluate the model.

```

35/35 31s 836ms/step

In [24]: # EVALUATION METRICS

```

accuracy = accuracy_score(y_true, y_pred)
precision = precision_score(y_true, y_pred, average='weighted')
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')

print("Accuracy:", accuracy)
print("Precision:", precision)

```

```
print("Recall:", recall)
print("F1-score:", f1)
```

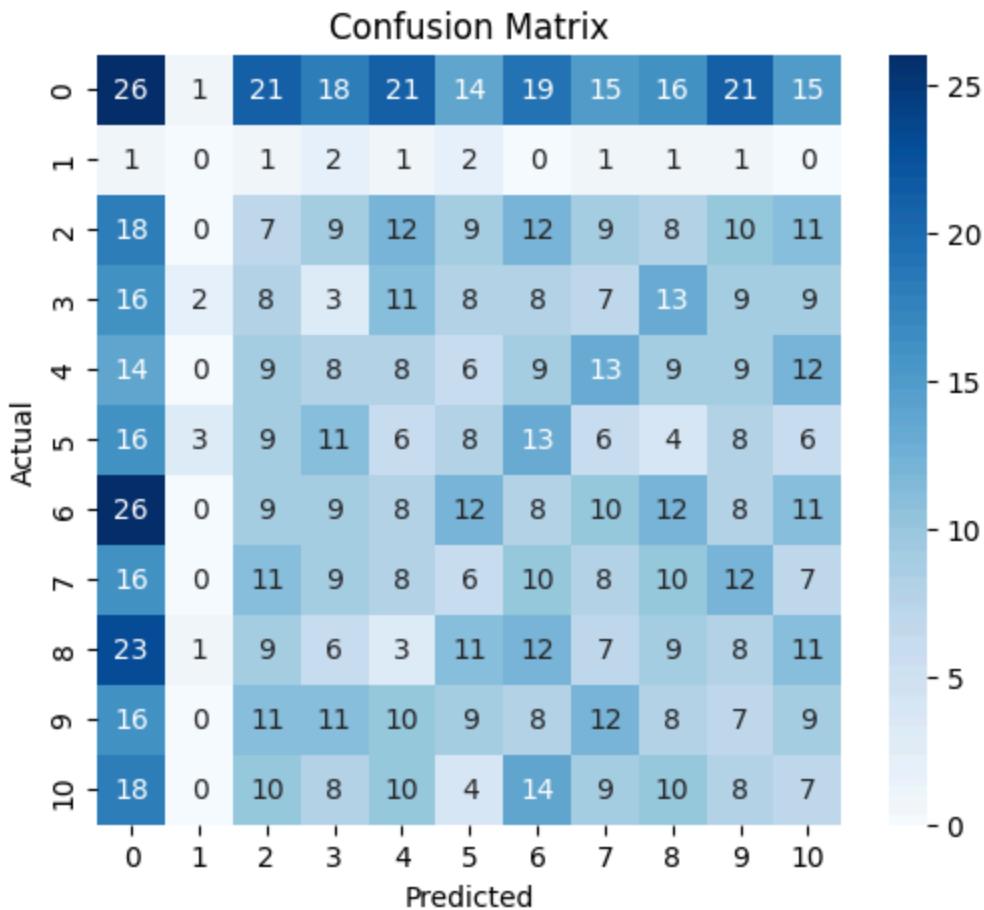
Accuracy: 0.0833333333333333
Precision: 0.08296495299596955
Recall: 0.0833333333333333
F1-score: 0.08314722577351684

In [25]: # CONFUSION MATRIX

```
cm = confusion_matrix(y_true, y_pred)
print(cm)

plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
[[26  1 21 18 21 14 19 15 16 21 15]
 [ 1  0  1  2  1  2  0  1  1  1  0]
 [18  0  7  9 12  9 12  9  8 10 11]
 [16  2  8  3 11  8  8  7 13  9  9]
 [14  0  9  8  8  6  9 13  9  9 12]
 [16  3  9 11  6  8 13  6  4  8  6]
 [26  0  9  9  8 12  8 10 12  8 11]
 [16  0 11  9  8  6 10  8 10 12  7]
 [23  1  9  6  3 11 12  7  9  8 11]
 [16  0 11 11 10  9  8 12  8  7  9]
 [18  0 10  8 10  4 14  9 10  8  7]]
```



In [26]: # STORE METRICS

```
results = []

results.append({
    "Model": "CNN_Model_1",
    "Accuracy": accuracy,
    "Precision": precision,
    "Recall": recall,
    "F1-score": f1
})

results_df = pd.DataFrame(results)
print(results_df)
```

	Model	Accuracy	Precision	Recall	F1-score
0	CNN_Model_1	0.083333	0.082965	0.083333	0.083147

In [27]: # VISUALIZE TRAINING HISTORY (ACCURACY AND LOSS)

```
def plot_training_history(history, model_name):
    plt.figure(figsize=(12, 4))

    # ----- ACCURACY -----
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{model_name} - Accuracy')
```

```

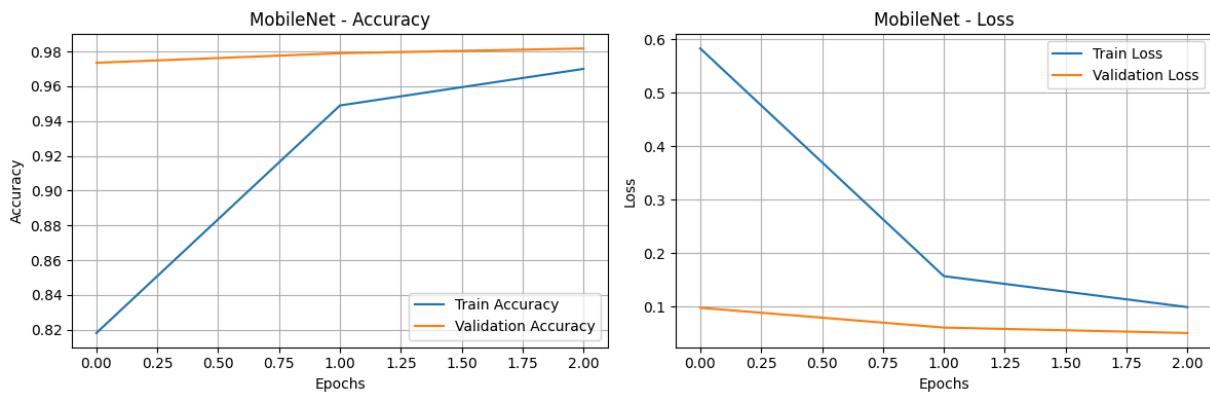
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

# ----- LOSS -----
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{model_name} - Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

plot_training_history(history_mobilenet, "MobileNet")

```



In [28]: # SAVE THE MODEL

```
best_model.save("best_mobilenet.h5")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

In [29]: # SAVE CLASS NAMES

```

import json

class_names = list(train_generator.class_indices.keys())

with open("class_names.json", "w") as f:
    json.dump(class_names, f)

print("Saved class_names.json")

```

Saved class_names.json

In [32]: # TEST SCALING AND GENERATORS

```
from tensorflow.keras.applications.mobilenet import preprocess_input

test_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input
)

test_generator = test_datagen.flow_from_directory(
    r"D:\DS PROJECTS SUBMISSION\Multiclass_Fish_Image_Classification\Dataset\images",
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    shuffle=False
)
```

Found 3187 images belonging to 11 classes.

In [33]: # FINAL TEST SCORE

```
test_loss, test_accuracy = best_model.evaluate(test_generator)
print("Test Accuracy:", test_accuracy)
print("Test Loss:", test_loss)
```

```
100/100 ━━━━━━━━━━ 103s 990ms/step - accuracy: 0.9906 - loss: 0.0244
Test Accuracy: 0.9905867576599121
Test Loss: 0.024425452575087547
```

In []: